

# Jet-Based Joint Embedding Predictive Architecture (J-JEPA)

CCNSB Seminar

**Abhiram Tilak**

*CND Dual Degree  
IIIT Hyderabad*

January 22, 2025

# Table of Contents

- 1 Introduction
- 2 Background
- 3 Architecture
- 4 Datasets and Training
- 5 Evaluation
- 6 Results
- 7 Conclusion

# Introduction

## About the Paper

This paper was showcased as a part of workshop ML4JETS2024 held on 4th November ([Link](#)). The paper was later published on 5th December 2024 to an Machine Learning conference called NeurIPS 2024.

Additional Information about the paper:

- This paper was an effort to adapt a well known Computer Vision Model called I-JEPA ([Image-based Joint-Embedding Predictive Architecture](#)) developed by Facebook Research in early 2023.

# Introduction

## About the Paper

This paper was showcased as a part of workshop ML4JETS2024 held on 4th November ([Link](#)). The paper was later published on 5th December 2024 to an Machine Learning conference called NeurIPS 2024.

## Additional Information about the paper:

- This paper was an effort to adapt a well known Computer Vision Model called I-JEPA ([Image-based Joint-Embedding Predictive Architecture](#)) developed by Facebook Research in early 2023.
- This paper serves as concurrent work for the model our lab has been working on since last semester.

# Background

## Motivations

### Problem with Supervised Models:

- Supervised learning typically acquire limited domain representations and focuses on a few key features for high prediction accuracy that must be learned anew for each task.

**Self-Supervised Learning:** A type of machine learning where models learn useful features and representations from unlabeled data  
SSL aims to learn generic representations summarizing domain features that prove useful across various downstream tasks. SSL tasks can be formulated on unlabeled data.

# Background

## Motivations

### Problem with Supervised Models:

- Supervised learning typically acquire limited domain representations and focuses on a few key features for high prediction accuracy that must be learned anew for each task.
- The problem comes when we have to scale the data being fed to the model.

**Self-Supervised Learning:** A type of machine learning where models learn useful features and representations from unlabeled data  
SSL aims to learn generic representations summarizing domain features that prove useful across various downstream tasks. SSL tasks can be formulated on unlabeled data.

# Background

## Motivations

### Problem with Supervised Models:

- Supervised learning typically acquire limited domain representations and focuses on a few key features for high prediction accuracy that must be learned anew for each task.
- The problem comes when we have to scale the data being fed to the model.
- A significant drawback of this is that the performance of ML models trained on simulations may not translate to real data, especially due to mismodeling in the former.

**Self-Supervised Learning:** A type of machine learning where models learn useful features and representations from unlabeled data  
SSL aims to learn generic representations summarizing domain features that prove useful across various downstream tasks. SSL tasks can be formulated on unlabeled data.

# Background

## Motivations

### Problem with Supervised Models:

- Supervised learning typically acquire limited domain representations and focuses on a few key features for high prediction accuracy that must be learned anew for each task.
- The problem comes when we have to scale the data being fed to the model.
- A significant drawback of this is that the performance of ML models trained on simulations may not translate to real data, especially due to mismodeling in the former.
- There have been efforts to run supervised models on huge datasets like ParT and OmniLearn-Alpha.

**Self-Supervised Learning:** A type of machine learning where models learn useful features and representations from unlabeled data  
SSL aims to learn generic representations summarizing domain features that prove useful across various downstream tasks. SSL tasks can be formulated on unlabeled data.



# Background

## Motivations for JEPA

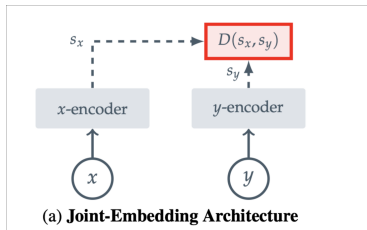
### Contrastive Methods (Also called Joint Embedding Architecture):

Contrastive Learning is a machine learning technique that teaches computers to understand similarities and differences by comparing pairs of data points.

It requires us to make several Data augmentations to the inputs, before we try to compare the difference between them.

These pretraining methods can produce representations of a high semantic level, but they also introduce strong biases that may be detrimental for certain downstream tasks or even for pretraining tasks with different data distributio

There have been a bunch of models in 2024 that use Contrastive methods like JetCLR, AnomalyCLR, BlackCLR, RS3L etc.



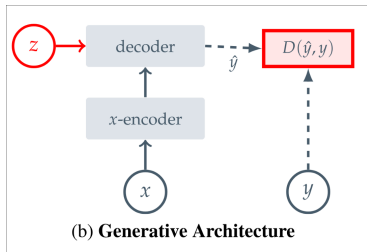
# Background

## Motivations for JEPA

**Generative Methods:** This idea is at the core of self-supervised generative methods, which remove or corrupt portions of the input and learn to predict the corrupted content. In particular, mask-denoising approaches learn representations by reconstructing randomly masked patches from an input, either at the pixel or token level.

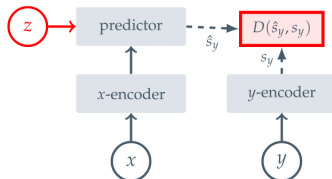
Masked pretraining tasks require less prior knowledge than view-invariance approaches and easily generalize beyond the image modality.

Only problem here is that they tend to underperform because they build only lower level semantic relationships between jet-representations.



# Architectre

## JEPA



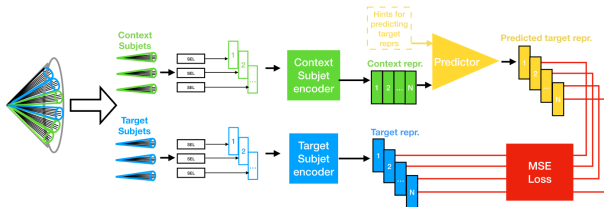
(c) **Joint-Embedding Predictive Architecture**

Given a jet, we recluster it into subjects, masking some as “target” subjects and defining others as “context” subjects. Then, we train a model to predict the representations of target subjects based on the representations of context subjects, using the positions of the target subjects as joint information.

This is also called Weakly-supervised model since the prediction stage uses some representations like positions of the target subjects and help the prediction.

# Architecture

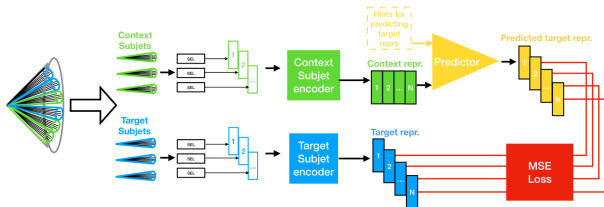
## J-JEPA



- The J-JEPA architecture begins by splitting the large-radius jet (large black cone) into target subjects and context subjects.

# Architecture

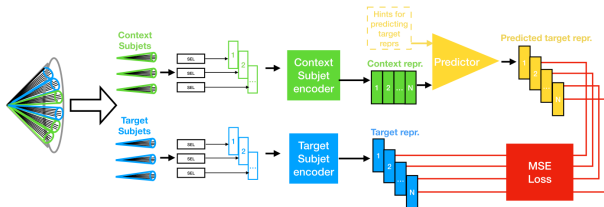
## J-JEPA



- The J-JEPA architecture begins by splitting the large-radius jet (large black cone) into target subjects and context subjects.
- The context encoder and target encoder then separately generate representations for the context subjects and the target subjects

# Architecture

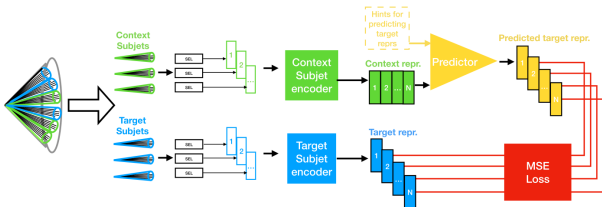
## J-JEPA



- The J-JEPA architecture begins by splitting the large-radius jet (large black cone) into target subjects and context subjects.
- The context encoder and target encoder then separately generate representations for the context subjects and the target subjects
- Using the positions of the target subjects as additional information (hints), the predictor takes the context representations and predicts the representations of the target subjects.

# Architecture

## J-JEPA



- The J-JEPA architecture begins by splitting the large-radius jet (large black cone) into target subjects and context subjects.
- The context encoder and target encoder then separately generate representations for the context subjects and the target subjects
- Using the positions of the target subjects as additional information (hints), the predictor takes the context representations and predicts the representations of the target subjects.
- Finally, the L2 loss function is used to compare the predicted target subject representations with the encoded target subject representations, minimizing the difference between them

# Datasets and Training

## Datasets:

**Pretraining:** JetClass: The pretraining task uses 1 Million jets, which is 1% of Jetclass dataset. It consists of 500k Top jets and 500k QCD jets.

**Finetuning:** TopTagging Reference Dataset: There are two scenarios in fine-tuning in this model. Both of them use TopTagging Dataset. The full dataset consists for 1.2 M Top jets.

Another scenario is using 10% of TopTagging dataset (120k jets). This is to represent situations where labeled training samples are limited.

**Training:** Experiments were performed on a single NVIDIA A100 GPU. For

- AdamW optimizer with cosine decay



# Datasets and Training

## Datasets:

**Pretraining:** JetClass: The pretraining task uses 1 Million jets, which is 1% of Jetclass dataset. It consists of 500k Top jets and 500k QCD jets.

**Finetuning:** TopTagging Reference Dataset: There are two scenarios in fine-tuning in this model. Both of them use TopTagging Dataset. The full dataset consists for 1.2 M Top jets.

Another scenario is using 10% of TopTagging dataset (120k jets). This is to represent situations where labeled training samples are limited.

**Training:** Experiments were performed on a single NVIDIA A100 GPU. For

- AdamW optimizer with cosine decay
- learning rate of  $10^{-3}$ , and weight decay  $10^{-2}$

# Datasets and Training

## Datasets:

**Pretraining:** JetClass: The pretraining task uses 1 Million jets, which is 1% of Jetclass dataset. It consists of 500k Top jets and 500k QCD jets.

**Finetuning:** TopTagging Reference Dataset: There are two scenarios in fine-tuning in this model. Both of them use TopTagging Dataset. The full dataset consists for 1.2 M Top jets.

Another scenario is using 10% of TopTagging dataset (120k jets). This is to represent situations where labeled training samples are limited.

**Training:** Experiments were performed on a single NVIDIA A100 GPU. For

- AdamW optimizer with cosine decay
- learning rate of  $10^{-3}$ , and weight decay  $10^{-2}$
- Epochs: 80, batch-size: 64

# Evaluation

## Metrics:

- Compare finetuned model, with model trained from scratch (no pretraining).

## Other Evaluations:

# Evaluation

## Metrics:

- Compare finetuned model, with model trained from scratch (no pretraining).
- Prediction Accuracy (number of correct classifications vs wrong)

## Other Evaluations:

# Evaluation

## Metrics:

- Compare finetuned model, with model trained from scratch (no pretraining).
- Prediction Accuracy (number of correct classifications vs wrong)
- Background rejection at signal efficiency 50%.

## Other Evaluations:

# Evaluation

## Metrics:

- Compare finetuned model, with model trained from scratch (no pretraining).
- Prediction Accuracy (number of correct classifications vs wrong)
- Background rejection at signal efficiency 50%.

## Other Evaluations:

- Using Traditional Embeddings vs. Custom Attention based embeddings

# Evaluation

## Metrics:

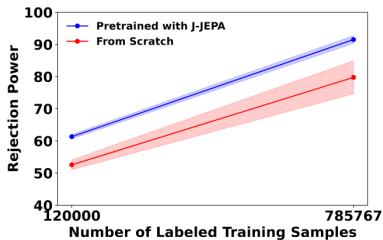
- Compare finetuned model, with model trained from scratch (no pretraining).
- Prediction Accuracy (number of correct classifications vs wrong)
- Background rejection at signal efficiency 50%.

## Other Evaluations:

- Using Traditional Embeddings vs. Custom Attention based embeddings
- Flattened subjects vs. Using class attention blocks to aggregate subject representations.

# Results

Model	Aggregation	Baseline 10%	Baseline Full	Finetuned 10%	Finetuned Full
Accuracy [%]					
SjT-T	Flatten	$87.52 \pm 0.16$	$89.13 \pm 0.10$	$88.21 \pm 0.55$	$89.95 \pm 0.13$
SjT-T	Cls Attn	$88.30 \pm 0.18$	$89.67 \pm 0.13$	$88.67 \pm 0.02$	$90.00 \pm 0.07$
AE-SjT-T	Flatten	$88.92 \pm 0.15$	$90.01 \pm 0.08$	<b><math>88.94 \pm 0.13</math></b>	<b><math>90.03 \pm 0.07</math></b>
AE-SjT-T	Cls Attn	$88.84 \pm 0.21$	<b><math>90.03 \pm 0.05</math></b>	$88.82 \pm 0.11$	$90.00 \pm 0.12$
$1/\varepsilon_B(\varepsilon_S = 0.5)$					
SjT-T	Flatten	$40.50 \pm 1.26$	$70.70 \pm 1.46$	$53.67 \pm 9.97$	$90.06 \pm 3.80$
SjT-T	Cls Attn	$52.56 \pm 1.54$	$79.75 \pm 5.12$	$61.32 \pm 0.66$	$91.51 \pm 1.20$
AE-SjT-T	Flatten	$67.34 \pm 1.40$	$97.79 \pm 3.90$	<b><math>70.47 \pm 1.09</math></b>	$97.52 \pm 1.71$
AE-SjT-T	Cls Attn	$67.19 \pm 1.54$	<b><math>99.38 \pm 2.80</math></b>	$68.25 \pm 1.64$	$95.47 \pm 1.83$





# Conclusion

This paper will likely have a follow-up which scales the data up and uses different encoders and uses various different datasets.

The future work involves:

- Implementing physics-informed architectures for the context and target encoders, such as the Particle Transformer

# Conclusion

This paper will likely have a follow-up which scales the data up and uses different encoders and uses various different datasets.

The future work involves:

- Implementing physics-informed architectures for the context and target encoders, such as the Particle Transformer
- Alternative strategies for embedding and defining targets and context

# Conclusion

This paper will likely have a follow-up which scales the data up and uses different encoders and uses various different datasets.

The future work involves:

- Implementing physics-informed architectures for the context and target encoders, such as the Particle Transformer
- Alternative strategies for embedding and defining targets and context
- Generalize the JEPA scheme to different physics objects: particles, events, detector readout, etc.

# Thank You